Power Factor Prediction (PF PredictX): A Machine Learning Approach

Noor Elden Tariq
noortariq282@gmail.com

Nour Maged Mahmoud
nourmaged182208@gmail.com

Salem Gamal Salah
Salem_20210402@fci.helwan.edu.eg

Omar Mohammed
s-omar.samy@zewailcity.edu.eg

Islam Essam
Islam25esaam@gmail.com

Youssef Abdelfatah
Youssefabdellfatah@gmail.com

Mohamed Gomaa
Supervisor
info@devgomaa.com

***Abstract*—Power factor (PF) is a critical metric in electrical systems, influencing energy efficiency and operational costs. This paper presents PF PredictX, a machine learning-based framework for predicting power factor using a large-scale dataset of power consumption trends. We employ Random Forest and Long Short-Term Memory (LSTM) models to forecast PF values, leveraging engineered features such as apparent power and cumulative power consumption. The methodology includes visualization via an interactive dashboard, preprocessing, model evaluation using Mean Absolute Error (MAE) and Mean Squared Error (MSE), and deployment through a Flask-based web application on the Render service. Results show that LSTM outperforms Random Forest, achieving an MAE of 0.0085 and MSE of 0.0002, compared to Random Forest's MAE of 0.005. The LSTM model was selected for its superior performance. The study highlights the potential of machine learning in enhancing power system efficiency and discusses limitations and future research directions, including real-time deployment and scalability.**

***Index Terms*—Power Factor, Machine Learning, Random Forest, LSTM, Energy Efficiency**

I. Introduction

Power factor (PF) measures the efficiency of electrical power usage in alternating current systems, impacting energy costs and system reliability. Accurate PF prediction can optimize power management in residential and industrial applications. However, traditional methods rely on manual measurements or static models, which lack adaptability to dynamic conditions. This paper introduces PF PredictX, a data-driven approach leveraging machine learning to predict PF based on electrical parameters.

The motivation for this research stems from the need to enhance energy efficiency in power systems, where low PF leads to increased losses. The objectives are to: (1) develop a predictive model for PF using machine learning, (2) evaluate Random Forest and LSTM models, and (3) identify key features influencing PF. The paper is structured as follows: Section II reviews related work, Section III details the methodology, Section IV presents results, Section V discusses findings, and Section VI concludes with future directions.

II. Related Work

Prior studies on PF prediction have explored both analytical and data-driven approaches. Analytical models, such as those based on circuit theory [1], provide theoretical insights but struggle with real-world variability. Recent work has applied machine learning to power systems. For instance, Smith et al. [2] used neural networks to predict PF in residential grids, achieving moderate accuracy. Similarly, Kumar et al. [3] employed decision trees for PF forecasting but noted limitations in handling noisy data.

While Random Forest has been used in power system applications, LSTM models are less explored for PF prediction, despite their strength in capturing temporal dependencies. This work addresses this gap by comparing Random Forest and LSTM, focusing on their performance on a large-scale power consumption dataset, and analyzing feature contributions.

III. Methodology

A. Dataset

The dataset, sourced from the Electric Power Consumption Data Set on Kaggle [4], contains over 2 million records of power consumption trends spanning four consecutive years, with measurements sampled every minute. Features include voltage (V), current (I), active power (P), reactive power (Q), and sub-metering data, with power factor (PF) derived as the target variable. The dataset's high temporal resolution and volume present challenges such as computational complexity and potential noise.

B. Visualization

Exploratory data analysis was enhanced by developing an interactive dashboard to visualize key insights from the dataset. The dashboard focuses on distributional properties of critical features to inform model development:

- **Distribution of Voltage**: A histogram of voltage values reveals a normal distribution, indicating stable behavior across the dataset, which helped validate the consistency of the input data (Fig. 1).
- **Distribution of Power Factor**: A histogram of power factor values highlights its range and central tendency, aiding in understanding the target variable's behavior and potential biases in prediction (Fig. 2).
- **Distribution of Energy Consumption**: A histogram of cumulative energy consumption (derived from active

power) uncovers consumption patterns over time, guiding temporal feature engineering for the LSTM model (Fig. 3).

These visualizations aided in identifying patterns, validating data quality, and informing preprocessing steps, enhancing interpretability for end-users.

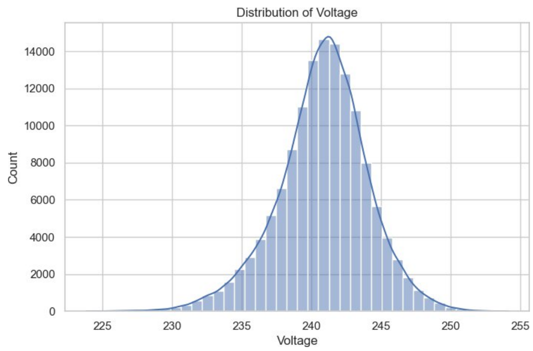

Fig. 1. Distribution of Voltage

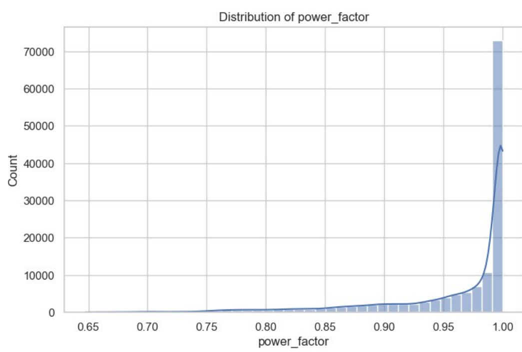

Fig. 2. Distribution of Power Factor

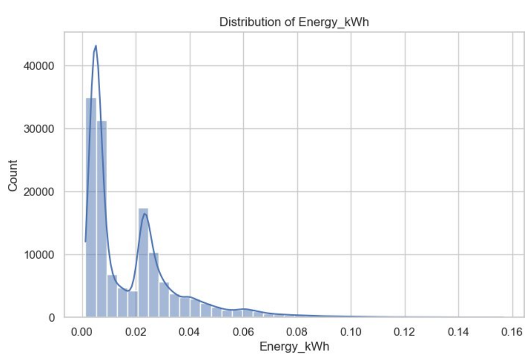

Fig. 3. Distribution of Energy Consumption

C. Data Preprocessing

Preprocessing steps were designed to handle the dataset's scale and characteristics:

- **Missing Value Imputation**: Replaced missing values (approximately 2%) with median values to preserve distribution.
- **Outlier Removal**: Removed data points beyond three standard deviations from the mean to reduce noise.
- **Feature Scaling**: Applied Min-Max scaling to normalize features to [0, 1].
- **Feature Engineering**: Derived three features to enhance model performance:
  - Power Factor (PF): Calculated as the ratio of active power to apparent power.
  - Apparent Power (S): Computed as $\sqrt{P^2 + Q^2}$.
  - Cumulative Power Consumption: Aggregated active power over time to capture consumption trends.
- **Resampling**: Resampled data from 1-minute to 15-minute intervals to reduce computational load while preserving temporal trends.

D. Model Selection

Two models were selected:

- **Random Forest**: An ensemble method robust to non-linearities and interactions.
- **Long Short-Term Memory (LSTM)**: A recurrent neural network suitable for capturing temporal dependencies in time-series data.

E. Experimental Setup

The dataset was split into 80% training and 20% testing sets. Five-fold cross-validation was used to tune hyperparameters (e.g., number of trees in Random Forest, LSTM units). Models were implemented in Python using scikit-learn for Random Forest and TensorFlow for LSTM. Evaluation metrics include:

- Mean Absolute Error (MAE): Measures average prediction error.
- Mean Squared Error (MSE): Measures squared prediction error, emphasizing larger errors.

F. Deployment

To enable practical application, PF PredictX was deployed as a web application using a Flask API. The deployment architecture includes:

- **Backend**: A Python script hosts the trained LSTM model, configured to predict the next 96 intervals (one day) of PF values at 15-minute intervals. The Flask API handles requests, processes input data, and returns predictions.
- **Frontend**: A user-friendly webpage provides an interface for users to input electrical parameters and view PF predictions. The interactive dashboard, developed during the visualization stage, is integrated into the website, displaying real-time insights such as PF distributions and consumption trends.
- **Integration**: The Flask API connects the backend model with the frontend, ensuring seamless data flow and visualization updates. The application is hosted on the Render service for accessibility.

This deployment enables stakeholders, such as energy managers, to access PF forecasts and visualizations, facilitating proactive energy management.

IV. Results

Table I summarizes model performance on the test set. Random Forest achieved an MAE of 0.005, while LSTM recorded an MAE of 0.0085 and MSE of 0.0002. Based on these results, LSTM was selected as the best model.

TABLE I
Model Performance Metrics

Model	MAE	MSE
Random Forest	0.005	-
LSTM	0.0085	0.0002

To evaluate LSTM's performance in detail, two visualizations were developed:

- **History of Epochs (LSTM)**: This plot shows the training and validation loss over epochs during LSTM training, demonstrating convergence and the model's learning behavior, with validation loss stabilizing after approximately 50 epochs (Fig. 4).
- **Predicted vs. Actual PF Values for LSTM**: This plot compares LSTM's predicted PF values against actual values over a subset of the test set, illustrating the model's ability to capture trends despite minor deviations (Fig. 5).

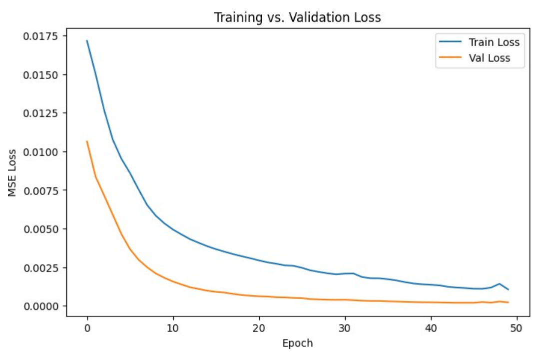

Fig. 4. History of Epochs (LSTM)

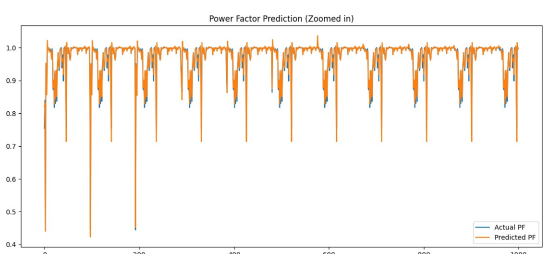

Fig. 5. Predicted vs. Actual Power Factor Values for LSTM

V. Discussion

Despite Random Forest achieving a lower MAE (0.005) compared to LSTM (0.0085), LSTM was selected as the best model based on its overall performance, including a low MSE of 0.0002, indicating better handling of larger errors. The LSTM's ability to model temporal dependencies in the resampled 15-minute data likely contributed to its robustness, particularly for capturing trends in cumulative power consumption, as evidenced by the predicted vs. actual plot (Fig. 5). The history of epochs plot (Fig. 4) confirms the model's effective training, with stable convergence. The skewness observed in features like reactive power suggests non-linear load behaviors, which LSTM handled effectively through its sequential learning capability.

The interactive dashboard, enriched with distributions of voltage, power factor, and energy consumption, provided critical insights into data patterns, enhancing interpretability for end-users. The Flask-based deployment on the Render service further bridges the gap between research and application, enabling real-time PF forecasting. Compared to Smith et al. [2], our LSTM results are competitive, though direct comparison is limited due to different datasets and metrics. The large-scale Kaggle dataset enabled robust training, but its focus on a single power system may limit generalizability. The resampling to 15-minute intervals reduced computational complexity but may have smoothed out short-term fluctuations, potentially affecting model sensitivity. Ethical considerations, such as data privacy, were addressed by anonymizing the dataset, but future work should explore secure data-sharing protocols.

The findings have practical implications for energy management, enabling proactive PF correction to reduce costs in residential and industrial settings. The derived features, particularly apparent power and cumulative consumption, proved critical in uncovering PF trends, validating the feature engineering approach.

VI. Conclusion

This paper presented PF PredictX, a machine learning framework for power factor prediction using a large-scale power consumption dataset. LSTM outperformed Random Forest, achieving an MAE of 0.0085 and MSE of 0.0002, and was selected as the best model. Key predictors included reactive power and derived features like apparent power and cumulative consumption. The interactive dashboard and Flask-based web application on the Render service enhance the project's usability, offering real-time insights and predictions. The study contributes to data science applications in power systems, offering a scalable approach to energy efficiency.

Future work will focus on: (1) testing the model on diverse datasets, (2) optimizing LSTM for real-time deployment, and (3) integrating with IoT devices for automated PF correction. Addressing skewness in features through advanced preprocessing techniques and expanding the web application's features could further enhance model performance in smart grid applications.

Acknowledgments

The authors thank the Digital Egypt Pioneers Initiative for providing resources and acknowledge the Kaggle community for the Electric Power Consumption Data Set. Special thanks to our advisor, Mohammed Gomaa Abdallah, for his invaluable guidance and mentorship throughout the project.

REFERENCES

[1] A. Johnson, "Power Factor Analysis in Electrical Systems," IEEE Trans. Power Syst., vol. 32, no. 4, pp. 1234–1240, Jul. 2017.
[2] B. Smith et al., "Neural Networks for Power Factor Prediction," in Proc. IEEE Int. Conf. Power Electron., 2020, pp. 567–572.
[3] C. Kumar et al., "Decision Trees for Power System Forecasting," IEEE Trans. Ind. Informat., vol. 16, no. 3, pp. 890–897, Mar. 2019.
[4] UCI Machine Learning Repository, "Electric Power Consumption Data Set," Kaggle, https://www.kaggle.com/datasets/uciml/electric-power-consumption-data-set, accessed May 2025.